

Feedback — IX. Neural Networks: Learning

[Help](#)

You submitted this quiz on **Sun 20 Apr 2014 10:36 PM PDT**. You got a score of **5.00** out of **5.00**.

Question 1

You are training a three layer neural network and would like to use backpropagation to compute the gradient of the cost function. In the backpropagation algorithm, one of the steps is to update $\Delta_{ij}^{(2)} := \Delta_{ij}^{(2)} + \delta_i^{(3)} * (a^{(2)})_j$ for every i, j . Which of the following is a correct vectorization of this step?

Your Answer

Score Explanation



$$\Delta^{(2)} := \Delta^{(2)} + \delta^{(3)} * (a^{(3)})^T$$



$$\Delta^{(2)} := \Delta^{(2)} + (a^{(3)})^T * \delta^{(3)}$$



$$\Delta^{(2)} := \Delta^{(2)} + \delta^{(3)} * (a^{(2)})^T$$



1.00

This version is correct, as it takes the "outer product" of the two vectors $\delta^{(3)}$ and $a^{(2)}$ which is a matrix such that the (i, j) -th entry is $\delta_i^{(3)} * (a^{(2)})_j$ as desired.



$$\Delta^{(2)} := \Delta^{(2)} + \delta^{(2)} * (a^{(3)})^T$$

Total

1.00 /

1.00

Question 2

Suppose `Theta1` is a 2x5 matrix, and `Theta2` is a 3x6 matrix. You set `thetaVec =`

[Theta1(:) ; Theta2(:)]. Which of the following correctly recovers Theta2?

Your Answer	Score	Explanation
<input type="radio"/> reshape(thetaVec(9:26), 3, 6)		
<input type="radio"/> reshape(thetaVec(10:27), 3, 6)		
<input type="radio"/> reshape(thetaVec(10:28), 3, 6)		
<input checked="" type="radio"/> reshape(thetaVec(11:28), 3, 6)	✓ 1.00	This choice is correct, since Theta1 has 10 elements, so Theta2 begins at index 11 and ends at index 11 + 18 - 1 = 28.
Total	1.00 / 1.00	

Question 3

Let $J(\theta) = 3\theta^3 + 2$. Let $\theta = 1$, and $\epsilon = 0.01$. Use the formula $\frac{J(\theta+\epsilon) - J(\theta-\epsilon)}{2\epsilon}$ to numerically compute an approximation to the derivative at $\theta = 1$. What value do you get? (When $\theta = 1$, the true/exact derivative is $\frac{dJ(\theta)}{d\theta} = 9$.)

Your Answer	Score	Explanation
<input checked="" type="radio"/> 9.0003	✓ 1.00	We compute $\frac{(3(1.01)^3 + 2) - (3(0.99)^3 + 2)}{2(0.01)} = 9.0003$.
<input type="radio"/> 8.9997		
<input type="radio"/> 11		
<input type="radio"/> 9		
Total	1.00 / 1.00	

Question 4

Which of the following statements are true? Check all that apply.

Your Answer	Score	Explanation
<input type="checkbox"/> Using a large value of λ cannot hurt the performance of your neural network; the only reason we do not set λ to be too large is to avoid numerical problems.	✓ 0.25	A large value of λ can be quite detrimental. If you set it too high, then the network will be underfit to the training data and give poor predictions on both training data and new, unseen test data.
<input checked="" type="checkbox"/> For computational efficiency, after we have performed gradient checking to verify that our backpropagation code is correct, we usually disable gradient checking before using backpropagation to train the network.	✓ 0.25	Checking the gradient numerically is a debugging tool: it helps ensure a correct implementation, but it is too slow to use as a method for actually computing gradients.
<input checked="" type="checkbox"/> If our neural network overfits the training set, one reasonable step to take is to increase the regularization parameter λ .	✓ 0.25	Just as with logistic regression, a large value of λ will penalize large parameter values, thereby reducing the changes of overfitting the training set.
<input type="checkbox"/> Gradient checking is useful if we are using one of the advanced optimization methods (such as in fminunc) as our optimization algorithm. However, it serves little purpose if we are using gradient	✓ 0.25	Gradient descent depends on the computation of correct gradient values at different parameter settings. Gradient checking ensures the computed values are correct.

descent.

Total	1.00 /
	1.00

Question 5

Which of the following statements are true? Check all that apply.

Your Answer	Score	Explanation
<input checked="" type="checkbox"/> Suppose you are training a neural network using gradient descent. Depending on your random initialization, your algorithm may converge to different local optima (i.e., if you run the algorithm twice with different random initializations, gradient descent may converge to two different solutions).	<input checked="" type="checkbox"/> 0.25	The cost function for a neural network is non-convex, so it may have multiple minima. Which minimum you find with gradient descent depends on the initialization.
<input type="checkbox"/> Suppose you have a three layer network with parameters $\Theta^{(1)}$ (controlling the function mapping from the inputs to the hidden units) and $\Theta^{(2)}$ (controlling the mapping from the hidden units to the outputs). If we set all the elements of $\Theta^{(1)}$ to be 0, and all the elements of $\Theta^{(2)}$ to be	<input checked="" type="checkbox"/> 0.25	Since the parameters are the same within layers, every unit in each layer will receive the same update during backpropagation. The result is that such an initialization does not break symmetry.

1, then this suffices for symmetry breaking, since the neurons are no longer all computing the same function of the input.

☒ Suppose we have a correct implementation of backpropagation, and are training a neural network using gradient descent. Suppose we plot $J(\Theta)$ as a function of the number of iterations, and find that it is **increasing** rather than decreasing. One possible cause of this is that the learning rate α is too large.

✓ 0.25

If the learning rate is too large, the cost function can diverge during gradient descent. Thus, you should select a smaller value of α .

☐ Suppose that the parameter $\Theta^{(1)}$ is a square matrix (meaning the number of rows equals the number of columns). If we replace $\Theta^{(1)}$ with its transpose $(\Theta^{(1)})^T$, then we have not changed the function that the network is computing.

✓ 0.25

$\Theta^{(1)}$ can be an arbitrary matrix, so when you compute $a^{(2)} = g(\Theta^{(1)} a^{(1)})$, replacing $\Theta^{(1)}$ with its transpose will compute a different value.

Total 1.00 / 1.00